

Interrupts and System Calls

Operating Systems I

Viktor Iakovlev (Victor Yacovlev)

x86 commands

Bytes	0	1	2	3	4	5
nop	0x00					
halt	0x10					
rrmovl <i>rA</i> , <i>rB</i>	0x20	<i>rA rB</i>				
irmovl <i>V</i> , <i>rB</i>	0x30	0x8 <i>rB</i>	<i>V</i>			
call <i>Dest</i>	0x80	<i>Dest</i>				
pushl <i>rA</i>	0xA0	<i>rA</i> 0x8				
popl <i>rA</i>	0xB0	<i>rA</i> 0x8				

Complex Instruction Set Computing

Commands are encoded using arbitrary byte size

Generic Purpose Registers

- Intel 8080/8085 - 8 bits (A,B,C,D)
- Intel 8086 - 16 bits (AX, BX, CX, DX)
- Intel 80386 - 32 bits (EAX, EBX, ECX, EDX)
- AMD Opteron - 64 bits
(RAX, RBX, RCX, RDX)

Intel 386 Architecture

- Builds for i386 are guaranteed to be working on any modern PC
- Multithreaded applications relies on next generation i486 due to CMPXCHG support
- i686 architecture features (Ubuntu minimum required):
 - CMOV - Conditional Move
 - SYSENTER/SYSCALL

gcc -march=ИМЯ_ПРОЦЕССОРА или native
<https://gcc.gnu.org/onlinedocs/gcc/x86-Options.html#x86-Options>

Additional Command Sets

- Vector Computing (MMX, SSE, AVX)
- Virtualization Extensions (VT-X or AMD-V)

cat /proc/cpuinfo

Floating Point Support

- Classic x86: distinct chip called co-processor x87
- Cons: the main CPU sleeps until FPU processes it's own command
- x87 works with it's own stack, changing data to CPU using main memory
- The x87 is a Standard for IA-32
- ... even it is not really used
- x86-64 completely dropped x87

Vector Arithmetics

- Special registers called XMM/YMM/ZMM
- Horizontal - maps vector to scalar
- Vertical - maps vector(s) to vector

*Remember on low-precision **float** type?
It might be useful!*

An Example: **dpps** command

dpps xmm0, xmm1, 0xF1

- Calculates a **dot product** for a vector of 16 floating-point values
- The constant only last operand describes:
 - High 4 bits (0xF) encodes what part of vectors to be in use
 - Low 4 bits (0x1) encodes where to place the result

Intel Command Sets


- MMX - integer only (128 bits)
- SSE - adds support for floating (128 бит)
- AVX и AVX-2 - floating point and integers (256 бит)
- AVX-512 - floating point and integers (512 бит)
Notice: only Hi-End processors supports it
- Some compilers provides support for intrinsics to be used with C/C++ language, but not ASM:

<https://software.intel.com/sites/landingpage/IntrinsicsGuide/>

Why to learn an assembly at «Operating Systems» course?

The Interrupts

Imaging a Hardware



Hardware Interrupts:

1. Button press
2. The next byte received by a transmission bus
3. Timer event
4. Aliens arrived

Interrupts Handling

- Save Instruction Pointer to stack
- Jump to instruction pointed by IDTR register + offset

What Happens on Interrupt

Before Interrupt Handle:

- Pipeline flush and disabling Out-of-order execution
- Virtual to physical mapping address table switch
- Switch stack pointer

To be Processed by Handler:


- Save processor state to stack
- [handle an event]
- Restore processor state back from stack

Event Handler

- Interact with hardware caused interrupt
- Requires direct access to hardware and memory
- Regular programs can't do it
*(even started by **root** superuser)*

x86 Isolation Rings

- Each process has it's own memory
- Regular programs has no access to I/O ports, but privileged has



Software Interrupts

- **int NN** for the most instruction sets (Intel, ARM etc.)
- Have the same behaviour like hardware interrupts
- Handled by BIOS before OS loaded
- Operating System might (but not required) to modify table of interrupt handlers

BIOS Interrupts Example

AH used for command, AL - for it's operand

BIOS

- INT 0x10 - Text output
- INT 0x13 - Disk I/O
- INT 0x15 - UART (COM-port) handling
- INT 0x16 - Keyboard

DOS

- INT 0x21 - used by DOS API

The Kernel

- Just a program that operates at highest (in most cases) processor level
- Has access to everything

System Call

- Some function from Kernel API
- Operates in Kernel Mode: the most privileged level
- Switches back to unprivileged level after execution finished

INT 0x80

- Unified interrupt number in Linux
- The EAX register stores the number of system call
(see `/usr/include/asm/unistd_32.h`)
- Parameters are stored at EBX, ECX, EDX, ESI, EDI, EBP
- The value returns within EAX register

Linux System Calls

- Documentation covered by section 2 of man pages
- The C standard library has C wrappers according to calling conventions

Example: string output

```
static const char S[] = "Hello";
write(1, S, sizeof(S));
/* man 2 write
    #include <unistd.h>
    ssize_t write(int fd, const void *buf, size_t count);
*/
```

```
#include <asm/unistd_32.h>
```

```
    mov eax, __NR_write // system call number
    mov ebx, 1           // the first argument
    mov ecx, S_ptr       // the second argument
    mov edx, 6           // the third argument
    int 0x80             // do it!
```

```
S:      .string "Hello"
S_ptr:  .dword S
```

linux-vdso.so (linux-gate.so)

```
$ ldd /usr/bin/cat
    linux-vdso.so.1 (0x00007ffe3ea87000)
    libc.so.6 => /lib64/libc.so.6 (0x00007f9cea333000)
    /lib64/ld-linux-x86-64.so.2 (0x00007f9cea6d6000)
```

- Virtual «library» to wrap some system calls as regular functions
- Maps functions that do not require high privileges:
 - `__vdso_clock_gettime`
 - `__vdso_getcpu`
 - `__vdso_gettimeofday`
 - `__vdso_time`

Modern Processors System Call Support

- `sysenter/sysexit` - for IA-32
- `syscall` и `vsyscall` - for AMD x86-64

