

Processes

Lesson 09 # Operating Systems I
Viktor Iakovlev (Victor Yacovlev)

The Process is...

- An instance of some running program
- An isolated address space of memory in the system

Process Attributes

Memory-Related Attributes:

- Processor Registers
- Memory Translation Tables
- Private and Shared memory pages
- Files mapping
- Individual system calls stack

Process Attributes

FS-Related Attributes:

- Descriptors table
- Current Working Directory
why there is no `cd` program?
- Корневой каталог
to be changed by superuser

Process Attributes

Other Attributes:

- Environment Variables
- Limits
- Resource Counters
- User and Group Identifiers

How to get an Information

- Command `ps` - show process list
- Command `top` - resources usage
- File System `/proc`

Жизненный цикл процесса

- **Running**
- **sTopped**
- **Suspended**
 - **Suspended** - can be terminated
 - **Disk Suspended** - can not be terminated
- **tracing**
- **Zombie**

State change examples

```
sleep(10);  
// from R to S
```

```
read(0, buffer, sizeof(buffer));  
// possible from R to S
```

```
read(fd, buffer, sizeof(buffer)); // in case of file  
// possible from R to D
```

```
_exit(5);  
// from R to Z
```

```
raise(SIGSTOP);  
// from R to T
```


Idleness

```
while (1) {  
    // do nothing - just waste CPU  
}
```

```
while (1) {  
    sched_yield(); // OK  
}
```

Process Creation

pid_t result = **fork()**

Creates a copy of current process

- `-1 == result` -- error
- `0 == result` -- for a child
- `0 < result` -- for a parent, in that case
result - is a Process ID for a newly created process

Example

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>

int main() {
    pid_t result = fork();
    if (-1 == result) { perror("fork :-("); exit(1); }
    if ( 0 == result) { printf("I'm son!\n"); }
    else {
        printf("I'm parent!\n");
        int status;
        waitpid(result, &status, 0);
        printf("Child exited with status %d\n", status);
    }
}
```

Process Copy

- Memory, registers etc. - an exact copy (except %eax/%rax)
- Differs in child process:
 - Process ID [getpid()], Parent ID [getppid()]
 - Pending Signals
 - Timers
 - File locks

Process Copy Side-Effects

```
int main() {  
    printf("abrakadabra ");  
    pid_t result = fork();  
    if (0==result) {  
        printf("I'm son\n");  
    }  
    else {  
        printf("I'm parent\n");  
    }  
}
```

abrakadabra I'm son
abrakadabra I'm parent

When it is not possible to start a process

- `/proc/sys/kernel/pid_max` [32768]
Maximum count of processes in the system
- `/proc/sys/kernel/threads-max` [91087]
Maximum count of threads, so each process have at least one thread

shell> :(){ :|:& };;:

Disclaimer!

This code is too dangerous!

```
void fork_bomb() {  
    pid_t p;  
    do {  
        p = fork();  
    } while (-1 != p);  
    while (1) sched_yield();  
}
```

Process Tree

- The first has number 1 - **init systemd**
- All but **init systemd** have a parent
- Process 1 becomes a parent of orphaned process in case of parent death
- Parents are noticed in case of child death

Process Termination

- System call `_exit(int)`
- Function `exit(int)`
- Operation return INT at main

```
printf("abrakadabra");  
_exit(0)
```

```
printf("abrakadabra");  
exit(0)
```

Process Termination

- Function `exit`:
 - calls all handlers registered by `atexit`
 - flushes I/O buffers
 - calls system call `_exit`

Exit Status

- Process might exit itself by calling system call `_exit(0≤code≤255)`
- Process might be externally terminated by someone using signals

```
int status;
waitpid(child, &status, 0);
if (WIFEXITED(status)) {
    printf("Exit code: %d", WEXITSTATUS(status));
}
else if (WIFSIGNALED(status)) {
    printf("Terminated by %d signal", WTERMSIG(status));
}
```

Zombie (<defunc>) Processes

- Zombie is the last process state
- Zombies are deleted from process table by parent which calls `wait` or `waitpid`
- Zombies must be cleared to prevent fork-bomb effect



exec - change current process to another program

man 3 exec

```
int execl(const char *path, const char *arg, ..., /* 0 */)
int execlp(const char *path, const char *arg, ..., /* 0 */)
int execlx(const char *path, const char *arg, ..., /* 0 */, char * envp[])

int execv(const char *path, char * const argv[])
int execvp(const char *path, char * const argv[])

#ifdef _GNU_SOURCE
int execvpe(const char *path, char * const argv[], char * const envp[])
#endif
```

Example

```
int main() {
    pid_t  pid = fork();
    if (-1==pid) { perror("fork :-("); exit(1); }
    if (0==pid) {
        execvp("ls", "ls", "-l", NULL);
        perror("exec :-(");
        exit(2);
    }
    else {
        waitpid(pid, NULL, 0);
    }
}
```

Example

```
int main() {
    pid_t  pid = fork();
    if (-1==pid) { perror("fork :-("); exit(1); }
    if (0==pid) {

        // a place to make additional process
        // setup: environment and others

        execlp("ls", "ls", "-l", NULL);
        perror("exec :-(");
        exit(2);
    }
    else {
        waitpid(pid, NULL, 0);
    }
}
```

Example

```
int main() {
    pid_t  pid = fork();
    if (-1==pid) { perror("fork :-("); exit(1); }
    if (0==pid) {
        chdir("/usr/bin");
        int fd = open("/tmp/out.txt",
                      O_WRONLY|O_CREAT|O_TRUNC, 0644);
        dup2(fd, 1); close(fd);
        execlp("ls", "ls", "-l", NULL);
        perror("exec :-(");
        exit(2);
    }
    else {
        waitpid(pid, NULL, 0);
    }
}
```