

# Signals:

## Interprocess Communication

Operating Systems I

Viktor Iakovlev (Victor Yacovlev)

# The Process Termination

- By process itself:
  - return from **main**
  - call function **exit**
  - system call **\_exit**
- Suitable for time-limited programs that have begin and end

# The Process Termination

- Killing by signal:
  - command **kill**
  - command **killall**
  - launched by **timeout**
  - keys Ctrl+C
  - terminal tab closed
  - system shutdown
- Asynchronous events to be processed by application

# The Process Termination

- Caused by error:
  - segmentation fault
  - broken pipe
  - division by zero
  - illegal instruction
  - assertion failed
- Default action is to kill process by sending a **signal**

# Signal

- Asynchronous event sent by:
  - the kernel
  - or by some other process
- Event handling:
  - just ignore it
  - terminate process
  - change state: s**T**opped | **R**unning
  - custom handling by application

# Signals

Number	Name	Default Action	Description
1	SIGHUP	Term	lost terminal connection
2	SIGINT	Term	Ctrl+C
3	SIGQUIT	Core	Ctrl+\
4	SIGILL	Core	illegal instruction
6	SIGABRT	Core	abort()

**man 7 signal**

*Демонстрации: 1) посылка сигналов; 2) do\_abort.c*

# Core Dump

- Process Memory Dump
- May be used for debugging purposes
- Managed by systemd in modern systems

```
/usr/sbin/sysctl kernel.core_pattern  
ulimit -c
```

Number	Name	Default Action	Description
9	SIGKILL	Term	process kill
11	SIGSEGV	Core	memory error
13	SIGPIPE	Term	broken pipe
15	SIGTERM	Term	process termination request
17	SIGCHLD	Ign	child terminated
18	SIGCONT	Cont	command fg
19	SIGSTOP	Stop	Ctrl+Z
23	SIGURG	Ign	socket urgent data

*Демонстрация: Ctrl+Z*



# Signals Handling [deprecated]

```
#include <signal.h>
```

```
// sighandler_t - Linux only  
typedef void (*sighandler_t)(int);
```

```
sighandler_t  
signal(int signum, sighandler_t handler);
```

# Standards and Systems

- System-V (Solaris)
- BSD and Linux

gcc -std=c99    v.s.    gcc -std=gnu99

~~#define \_BSD\_SOURCE~~

#define \_DEFAULT\_SOURCE

#define \_GNU\_SOURCE

# Signals Handling [deprecated]

```
#include <signal.h>
```

```
// sighandler_t - Linux only  
typedef void (*sighandler_t)(int);
```

```
// sig_t - BSD only  
typedef void (*sig_t)(int);
```

```
sighandler_t  
signal(int signum, sighandler_t handler);
```

# Signals Handling [modern-way]

```
#include <signal.h>
```

```
struct sigaction {  
    void      (*sa_handler)(int);  
    void      (*sa_sigaction)(int, siginfo_t *, void *);  
    sigset_t  sa_mask;  
    int       sa_flags;  
    void      (*sa_restorer)(void);  
};
```

```
int  
sigaction(int signum,  
          const struct sigaction *act,  
          struct sigaction *oldact /* might be NULL */);
```

# System call **signal**

- To be used for default signal handlers:  
SIG\_IGN (=1) и SIG\_DFL (=0)
- For Linux systems works mostly like System-V

# Signal Handler

- Might be called at **arbitrary time**:
  - uses the current stack
  - it is warranted for x86\_64 to have a «RedZone» by size of 128 bytes
  - Might use only restricted set of functions:  
Async-Signal Safe (AS-Safe)

# Async Signal Safety

- Functions safeties:
  - Unsafe
  - Multi-Threading Safe (MT-Safe)
  - Async-Signal Safe (AS-Safe)
- The sets of AS-Safe and MT-Safe are not equal! Example: **fwrite**
- See for a full list in POSIX:  
man 7 signal-safety

# Signal Handler

- The extra integer type **sig\_atomic\_t**
  - implemented as **int** for most platforms
  - guaranteed to be atomic at signal-level, but not MT-level
- The **volatile** language keyword
  - tells to compiler not to optimize the variable



# Examples: what to handle

- SIGTERM and SIGINT - terminate correctly
- SIGINT might ask user if he sure to terminate
- SIGHUP - reload settings file(s) if any
- SIGCHILD - read child exit status
- SIGSEGV - might try to save critical data

# Pending Signals

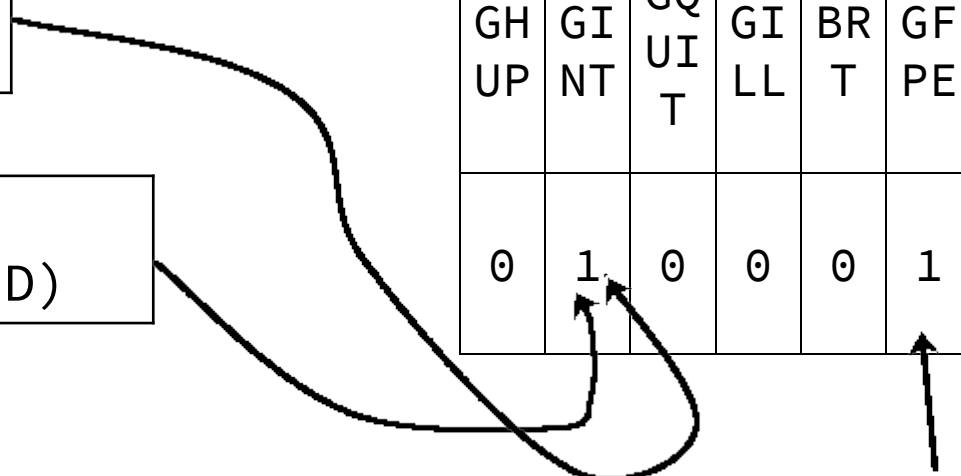
- The process attribute
- Not to be inherited by child while fork

**Process 1**  
`kill(SIGINT, PID)`

**Process 2**  
`kill(SIGINT, PID)`

**Process 3**  
`kill(SIGFPE, PID)`

SIGHUP	SIGINT	SIGQUIT	SIGILL	SIGABRT	SIGFPE	SIGKILL
0	1	0	0	0	1	0



# Signal Mask

Signals to be processed →

Mask of allowed/blocked signals →

0	0	0	0	0	1	0
1	0	1	1	1	1	1
SIGHUP	SIGINT	SIGQUIT	SIGILL	SIGABRT	SIGFPE	SIGKILL
0	1	0	0	0	1	0

**Process 1**

`kill(SIGINT, PID)`

**Process 2**

`kill(SIGINT, PID)`

**Process 3**

`kill(SIGFPE, PID)`

**Signal Mask,**  
*opposite to pending signals,*  
***to be inherited by childs while fork***

# POSIX Signals

- Have standard behaviour for most UNIX-like variants
- Processes are noticed by signal facts, but not their count
- Signals might be processed in random order

# Extended Signals

POSIX.1b real-time extensions; have been implemented in Linux

- Values from `SIGRTMIN` to `SIGRTMAX`
- Might be used as regular UNIX-signals (by system call `kill`)
- Might be queued and counted in specific order
- Might carry some trivial data

# Sending Signals Using Queue

## POSIX:

```
sigqueue(int pid,  
          int signum,  
          const union signal value)
```

```
union signal {  
    int    sival_int;  
    void*  sival_ptr;  
};
```

## LINUX:

```
sys_rt_sigqueueinfo(  
    int pid,  
    int signum,  
    const siginfo_t  
        *uinfo  
)
```

