

Inter-Process Communications.

Pipes

Operating Systems part II
Viktor Iakovlev (Victor Yacovlev)

UNIX Process

- Isolated virtual memory
- Implemented by CPU
- Side Effect: it is required to communicate via Kernel to make processes to interact

Process Interaction


- Signals
 - UNIX Signals
 - POSIX Real-Time Extensions
- Shared Memory Mapping:
mmap(..., MAP_SHARED, ...)
- Regular Files

UNIX Pipelines

	ls	-l	 	wc
#	command	argument	vert	command

- Two programs starts at a time
- Out of first one redirects to input of second one
- Implemented using **unnamed pipe**

UNIX Pipe




- System call `pipe(int fds[2])`
 - out-parameter - where to store fds
 - `fds[0]` - read-only, `fds[1]` - write-only

Pipes usage

- Interprocess Communications:
 - file descriptors are inherited by child processes
 - fork+exec keeps created pipe
- Within the same process:
 - one-threaded case: like a buffer
 - might be used for multithreading


UNIX Pipe



$$V = \pi R^2 \cdot L = 65356 \text{ bytes}$$

(old Linux Kernels: only 4096 bytes)

UNIX Pipe



$$Q = \pi R^2 \cdot B = \text{PIPE_BUF bytes}$$

- Atomicity is warranted for blocks of size PIPE_BUF
- For Linux (x86) PIPE_BUF value is 4096
- POSIX.1-2001 requires PIPE_BUF \geq 512

Pipe write

- Its **required** than reading part of pipe is opened by some process
- If no more space in pipe buffer:
 - blocks until free space; or
 - special mode `O_NONBLOCK` fails write system call with `EAGAIN`
- «Broken pipe» error occurs in case of no one readers exist

Non-Blocking Read/Write Operations

- `O_NONBLOCK` flag for open system call
- Reading and writing operations do not the process, but fails with `EAGAIN` error
- Might be controlled by `fcntl`

Pipe read

- If at least one byte in the buffer: no block, just read
- Waits until data ready (except `O_NONBLOCK` mode)
- If there is no writing side (e.g. opposite side of pipe was closed) - works like end-of-file

Dead Lock

- Trying to read a pipe, which has will be never written
- Usually due to forgotten `close` in some parent process

UNIX Pipelines

	ls	-l	 	wc
#	command	argument	vert	command

- 0 and 1 - stdin and stdout
- pipe(int fds[2]) created two **random** descriptor numbers

dup

- **dup2** - create a copy of file descriptor
- The copy is just a reference for file-like object

dup v.s. dup2

```
int dup2(int oldfd, int newfd);
```

1. Close newfd in case if in use
2. newfd = dup(oldfd)

Those two operations processed at once!

FIFO: Named Channels

- Unnamed channels might be used only by **relative** processes

Example:

```
ls -l | wc
```

bash process creates a pipe, then creates two childs and both of them inherits the pipe

- Named channels (FIFO) allows to interact non-relatice processes

FIFO: Named Channels

- FIFO - special file type
- mkfifo - the command and function
- The special file is opened like a regular file
- But behaviour like a pipe

Demo: 1) mkfifo channel && strace cat channel; 2) echo hello >channel

Pipes in Real Life

- File descriptors 0, 1 and 2
- Might communicate several programs via pipeline
- Other cases: launch program in Terminal or IDE
- Usually unnamed pipes are in use, but not FIFO

Pipe's Limitations

- Unidirectional: requires two pipes for bidirectional communications
- Synchronization is not easy when dealing with multiple readers/writers

Sockets v.s. Pipes

- Bidirectional
 - Allows to detect events, but pipes - not
 - Unified API for local and network communications
-
- Much more complex than pipes

The simplest socket: socket pair

- FreeBSD and Linux only, but not any UNIX
- Bidirectional channel for link

```
socketpair(AF_LOCAL,  
            SOCK_STREAM,  
            0,  
            out: int fds[2]  
            ) → err;
```

Sockets

- The topic of next week
- Many types of sockets are declared, but two types in use:
 - local to computer (AF_LOCAL)
 - network (AF_INET or AF_INET6)