

# Multithreading

Operating Systems II

Viktor Iakovlev (Victor Yacovlev)

# Threads

- Thread (aka lightweight process) - task to be controlled by scheduler
- Each task has it's own stack and processor state
- One process might have several threads

# Main Thread

- Function `_start` starts execution in main thread.
- Main thread might start other threads
- Any thread can start new thread
- All threads are equal

# exit

- System call `_exit` terminates the process: all threads
- System call `pthread_exit` terminates current thread only
- It is possible to terminate main thread but keep other threads working

# Stack

## Память процесса:

|                          |                             |
|--------------------------|-----------------------------|
| 0x7fffffffffff<br>– 8 Mb | <b>Стек</b><br>↓            |
| 0x03106000               | ↑<br><b>Куча</b>            |
|                          | <b>.bss</b>                 |
|                          | <b>.data</b>                |
|                          | <b>.rodata</b>              |
|                          | <b>.text</b>                |
| 0x00400000<br>0x00000000 | <b>Заполнено<br/>нулями</b> |

## Память потока:

|                          |                  |
|--------------------------|------------------|
| 0x7fffffffffff<br>– 8 Mb | <b>Стек</b><br>↓ |
|--------------------------|------------------|

|                           |                  |
|---------------------------|------------------|
| 0x7ffffff7fffff<br>– 8 Mb | <b>Стек</b><br>↓ |
|---------------------------|------------------|

|                           |                  |
|---------------------------|------------------|
| 0x7ffffefffffff<br>– 8 Mb | <b>Стек</b><br>↓ |
|---------------------------|------------------|

Stack areas might be mixed with dynamic library areas

# Stack

- Thread stack size has the same size as main thread by default
- POSIX allows to specify stack size for newly created threads
- Guard Page below the stack
- **No memory isolation: any thread is allowed to access any thread's stack memory**

# Process v.s. Thread

- Just a task from system scheduler's point of view
- System call `sched_yield` affects current thread; in case of single-threaded programs - Main thread
- Process count limit for Linux - limit for maximum number of threads but not processes

# Process v.s. Thread

## Process

- Complete isolation from each other
- Process crash has no effect to other processes
- Processes might have different privileges
- Inter-Process Communication is complex in general

## Thread

- Threads share the same memory address space
- Signal (in case of errors too) affect all threads
- All threads work within the same user
- Simple communication using common variables



# Thread and Process Attributes

## Common to All Threads

- Process ID
- Parent process ID
- Group ID/Session ID
- Associated Terminal
- UID/GID
- File Descriptors
- File Locks
- Signal Handlers
- umask, cwd, root dir
- Limits

## Distinct

- Thread ID
- Thread Stack Memory
- Signal Handler Stack
- Scheduling Priority
- Signal Mask  
pthread\_sigmask
- errno Value

# <errno.h>

- «Global Variable» behaviour
- Implemented by macro in modern systems

```
/* The error code set by various library  
functions.  */
```

```
extern int
```

```
    *__errno_location (void)
```

```
    __THROW __attribute_const__;
```

```
# define errno (*__errno_location ())
```

# Threading API

- POSIX (Linux, \*BSD, Mac)  
    <pthread.h>  
    pthread\_create(...)
- WinAPI (Windows)  
    <Windows.h>  
    CreateThread(...)

# WinAPI Threads

```
HANDLE CreateThread(  
    LPSECURITY_ATTRIBUTES    lpThreadAttributes,  
    SIZE_T                    dwStackSize,  
    LPTHREAD_START_ROUTINE    lpStartAddress,  
    LPVOID                    lpParameter,  
    DWORD                     dwCreationFlags,  
    LPDWORD                   lpThreadId  
);
```

Minimum stack size - 64Kb

# C++11 Threads

- Common interface for all platforms  
    `<thread>`  
    `std::thread(...)`
- Covers only common features

# C11 Threads

- Common interface for all platforms

`<threads.h>`

```
void call_once(  
    once_flag *flag,  
    void (*func)(void)  
);
```

- **At present there is no any implementation**

# Threading and fork

- All threads are equal → it is possible to call `fork()` from any thread
- Process creation implies memory copy, including stacks of all threads
- Newly created process **will have just one task**, that continues execution of thread that called `fork()`

# Threading and fork

- Might have locked resources → deadlock
- In case of using fork, it is required to call `pthread_atfork` to register a function that releases all possible resources



# Threading and fork + exec

- System call `exec` completely replaces address space by another program
- System call `fork` keeps only one task running
- Combination of `fork+exec` is safe in general

# Signal Mask

(each thread has it's own)

Signalas to be processed →

Mask of [blocked] signals →

|        |        |         |        |         |        |         |
|--------|--------|---------|--------|---------|--------|---------|
| 0      | 0      | 0       | 0      | 0       | 1      | 0       |
| 1      | 0      | 1       | 1      | 1       | 1      | 1       |
| SIGHUP | SIGINT | SIGQUIT | SIGILL | SIGABRT | SIGFPE | SIGKILL |
| 0      | 1      | 0       | 0      | 0       | 1      | 0       |

**Process 1**


kill(SIGINT, PID)

**Process 2**

kill(SIGINT, PID)

**Process 3**

kill(SIGFPE, PID)



# Signals Processing

- Signal Mask is a **process** but not thread attribute
- Signal handlers are available for all threads
- Signal handlers are to be executed by **arbitrary** thread
- It is possible to block signals handling for distinct threads (pthread\_sigmask)

# Thread Cancellation

- Avoid using this
- Two ways to perform:
  - default: thread might be canceled at cancellation point (mostly system calls)
  - hard way: force cancellation even on code like this:  
`while (1) {}.`  
Threads can't be canceled immediately.

# Thread Cancellation

- Function `pthread_cancel` "cancels" specified thread: sends a request to stop
- Thread might disallow these requests:  
`pthread_setcancelstate(  
    PTHREAD_CANCEL_DISABLE)`

# Thread Cancellation

- Two ways of cancellation:
  - at cancellation point (man 7 pthreads)
  - by sending a signal
- Thread might specify cancellation type  
`pthread_setcanceltype`

# Thread Cancellation

- On thread canceled:
  - semaphores and mutexes might be locked
  - unclosed file descriptors
  - lost memory pointers
  - just important work not complete
- Avoid thread cancellation!

# Thread-Local Storage

- Global variables are available for all threads
- Language keywords `_Thread_local` (C11) and `thread_local` (C++) make variables to be local to threads
- Values of TLS-variables are not isolated



# Global Variables

- May lead to data race
- Might be changed indirectly via standard functions
- Many functions have marked as Thread-Safe: they are implemented with locks (with drawback of performance loss)
- Example: `getc` v.s. `getc_unlocked`

# How to Avoid Data Race

- Use atomic values

`<stdatomic.h>` (C11)

```
typedef _Atomic _Bool atomic_bool;  
typedef _Atomic char atomic_char;  
typedef _Atomic signed char atomic_schar;  
typedef _Atomic unsigned char atomic_uchar;  
typedef _Atomic short atomic_short;  
typedef _Atomic unsigned short atomic_ushort;  
typedef _Atomic int atomic_int;
```

It is required to use atomic operations

- Use locks: semaphores and mutexes

# The Key Problem of Multithreading

- There are a lot of single-threaded programs accumulated due computes history
- The simplest way make them parallel-safe is to add locks
- The simplest way is not the most effective

# Example: Python

- Thread start is `pthread_create` + initialization of internal structure `PyThreadState`
- Each Python Thread is a system Thread  
but...
- All of them uses the single interpreter instance (implemented as `PyInterpreterState` structure)

**Global Interpreter Lock (GIL)**

# Example: Python. How to Avoid GIL

- Do not use Python :)
- Use multiprocessing instead of multithreading in case to make program really parallel

# In-Kernel Threads

- Goal: make Kernel Task to run in parallel
- All of them runs in Kernel address space
- Linux assigns virtual ProcessID to all in-kernel threads like processes
- Examples:
  - kswapd - manages swap
  - kworker - [in most cases] handles I/O operations

# Giant Kernel Lock

- Global lock for the whole Kernel
- Processes might block another processes during system calls
- Problem origins from legacy device drivers and subsystems
- The problem firstly eliminated for BSD at 2009 (FreeBSD 8.0)
- The problem eliminated for Linux at 2011 году (Kernel version 2.6.39)

# BeOS / HaikuOS

<https://www.haiku-os.org/>

- Operating System that have support for multiple-cores by design
- Each Kernel subsystem works in it's own thread

